

Signal Implementation in xv6 Kernel

Project Signal - Final Report

Jaouhara Zerhouni Khal

Harrold Tok Kwan Hang

Hok Layheng

June 3, 2025

1 Team Members

1. **Jaouhara Zerhouni Khal** — 12211456
2. **Harrold Tok Kwan Hang** — 12212025
3. **Hok Layheng** — 12210736

2 Executive Summary

This project involved implementing a signal handling mechanism in an XV6-based operating system kernel. The signal system supports asynchronous communication with user processes and mimics POSIX-style signal behavior, enabling processes to send, receive, and handle signals through custom-defined or default behaviors. This report is a continuation from Checkpoint 1, which now describes the support of new functions such as SIGKILL, signal across fork and exec, and bonus functions (SIGCHLD, SIGINFO and SIGALARM).

3 Project Design and Architecture

Our signal implementation is integrated modularly within the xv6 kernel to ensure clarity and maintainability. The key design components include:

- **Signal Table:** Each process maintains a signal table storing registered handlers, pending signals, and the signal mask.
- **Kernel Modifications:** Core kernel files (`syscall.c` and `trap.c` were modified to manage signal delivery, context saving, and system call support (e.g., `sigaction`, `sigprocmask`, `alarm`).
- **Signal Delivery Logic:** Signals are delivered lazily during user-kernel transitions (e.g., timer interrupt, syscall return) to avoid reentrancy issues. Delivery respects masking and handler type (default, ignore, user-defined).

- **User Context Handling:** User-mode context is saved/restored during signal entry/return, enabling transparent resumption post-handler execution.
- **Extension Support:** Architecture supports additional features like `siginfo_t`, `SIGCHLD`, and timer-based `SIGALRM` via modular extensions and syscall interfaces.

Function Descriptions

`int siginit(struct proc *p)`

Initializes the signal-related structures inside a process control block (PCB). It sets all handlers to default, masks to 0, and clears pending signals.

`int siginit_fork(struct proc *parent, struct proc *child)`

Copies the signal action and mask from the parent process to the child during fork. Clears the child's pending signals. For kernel test mode, it sets `SIGUSR2` as pending for testing.

`int siginit_exec(struct proc *p)`

Called during an `exec` system call to reset signal handlers to default. Retains ignored signals explicitly by re-setting them to `SIG_IGN`.

`int do_signal(void)`

Called during user-to-kernel transitions. Handles pending, unmasked signals by either performing the default action or invoking a user-defined handler via stack manipulation.

`int sys_sigaction(int signo, const sigaction_t *act, sigaction_t *oldact)`

Implements the `sigaction()` system call. Sets or retrieves the signal handler and associated mask for a particular signal. Prevents overriding `SIGKILL`.

`int sys_sigreturn(void)`

Called from user space after a signal handler returns. Restores the saved context (registers, PC, and signal mask) from the stack and decreases the signal handling depth.

`int sys_sigprocmask(int how, const sigset_t *set, sigset_t *oldset)`

Modifies or retrieves the signal mask of the current process depending on the `how` argument: `SIG_BLOCK`, `SIG_UNBLOCK`, or `SIG_SETMASK`. `SIGKILL` cannot be masked.

`int sys_sigpending(sigset_t *set)`

Retrieves the set of pending signals for the current process.

`int sys_sigkill(int pid, int signo, int code)`

Sends a signal to the process with the given PID. Populates the corresponding `siginfo_t` structure and wakes the process if it is sleeping.

`int sys_alarm(unsigned int seconds)`

Sets an alarm for the calling process to go off after `seconds`. When the timer expires, a `SIGALRM` is sent to the process. If a previous alarm was set, the remaining time is returned. Passing `seconds = 0` cancels any pending alarm.

`void check_alarm(void)`

Called periodically by the kernel clock tick handler. Checks if the alarm time for the current process has been reached. If so, it delivers a `SIGALRM` and clears the alarm.

Implementation of Bonus

Bonus Checkpoint 5.3.1 – `SIGALRM`

Goal: Enable support for the `SIGALRM` signal to allow processes to set alarms that deliver a signal after a specified duration.

Implementation Summary:

- Implemented `sys_alarm()` in `ksignal.c`, which schedules or cancels alarms.
- The alarm time is tracked using `p->alarm_ticks` and `p->alarm_interval`, based on the system's tick rate (100Hz).
- Alarms are checked periodically in `check_alarms()`, which queues `SIGALRM` signals and clears expired timers.
- `sys_sleep()` and `check_sleeping_processes()` were implemented to manage process suspension and resumption based on ticks.
- `handle_timer_interrupt()` triggers both alarm and sleep checking during each timer interrupt.

Testing Summary:

- Used `test_alarm`, `test_alarm_remaining` and `test_alarm_cancel` to verify correct `SIGALRM` behavior.
- Validated:
 - Accurate remaining time when setting new alarms.
 - Correct delivery and cancellation of `SIGALRM`.
 - Proper handling of sleep and wake-up events.
- Assertions and debug output confirmed tick values, alarm triggering, and cancellation correctness.

Bonus Checkpoint 5.3.2 – `siginfo_t`

Goal: Enable the kernel to construct and deliver a `siginfo_t` structure to user-level signal handlers, as required by the `SA_SIGINFO` convention.

Implementation Summary:

- Extended the signal delivery path to build a `siginfo_t` structure dynamically on the user stack before invoking the handler.
- Populated fields:

- `si_signo` – the signal number.
 - `si_code` – the reason or metadata for the signal.
 - `si_pid` – the sender's PID or `-1` if from the kernel.
- Updated `sys_sigkill()` to populate `siginfo` metadata used for constructing the structure.
- The main logic was implemented in:
 - `handle_stack()` – builds and pushes `siginfo_t` and `ucontext_t` onto the stack.
 - `do_signal()` – triggers signal delivery using `handle_stack()`.
 - `sys_sigreturn()` – restores the context and signal mask from the user stack.

Testing Summary:

- Verified using the `siginfo_bonus()` test case.
- A child process received:
 - A signal from a user process (`SIGUSR1`).
 - A simulated kernel signal (`SIGUSR2`).
- Inside the handler, assertions verified the accuracy of `si_signo`, `si_pid`, and `si_code`.
- The test passed only when both signals were handled with the correct metadata.

Bonus Checkpoint 5.3.3 – SIGCHLD

Goal: Automatically deliver `SIGCHLD` to a parent process when a child exits, including signal metadata.

Implementation Summary:

- Modified `exit()` in `proc.c` to check the parent's signal handler for `SIGCHLD`.
- If set (not `SIG_DFL` or `SIG_IGN`), `SIGCHLD` is delivered with:
 - `si_signo = SIGCHLD`
 - `si_pid = child PID`
 - `si_code = exit code`
- The parent is woken up if it is sleeping, to handle the signal promptly.

Testing Summary:

- The `sigchld_bonus()` test case verifies correct signal handling.
- A parent registers a `SIGCHLD` handler, and a child exits with code `123`.
- Inside the handler:
 - Assertions check `info->si_signo == SIGCHLD`.

- `wait()` is called to reclaim the child.
- Verified that `info->si_pid` and `info->si_code` match the expected values.
- The test succeeds only if all validations pass, ensuring correct `SIGCHLD` signal delivery and metadata.

4 Challenges and Solutions

4.1 Challenge 1: Integrating Stack Setup Inline

Problem: We initially factored stack setup into a separate `handle_stack()` helper, but inlining that logic into `do_signal()` risked code duplication and made the signal-delivery loop bulky. **Solution:** We carefully merged the stack-alignment, `ucontext siginfo_t` allocation and user-copy calls directly into `do_signal()`, eliminating the extra function call while retaining clear comments and minimizing duplicated lines by reusing a small block that saves registers and copies both structures under the same lock.

4.2 Challenge 2: Mask and Nesting Management

Problem: Ensuring that the original signal mask is preserved across nested handlers and correctly restored by `sys_sigreturn()`, while also preventing the handler itself from being interrupted by the same signal. **Solution:** We embed the old mask in the saved `ucontext.uc_sigmask` and, upon entry, OR in both the handler's `sa_mask` and the signal's own bit to the process `sigmask`. In `sys_sigreturn()`, we restore exactly the saved mask and decrement `handling_depth`, ensuring proper nesting behavior.

4.3 Challenge 3: Atomicity and Lock Ordering

Problem: Avoiding races between `sys_sigkill()` enqueueing new signals and `do_signal()` clearing them, especially when a process is sleeping or executing user-copy operations under the mm lock. **Solution:** In `sys_sigkill()`, we acquire the target's `p->lock` before touching `sigpending` and release only after waking it. In `do_signal()`, we disable interrupts around the pending-check and then acquire `mm->lock` briefly to copy out both `siginfo_t` and `ucontext` in one atomic section, avoiding inversion by always grabbing `p->lock` before `mm->lock`.

4.4 Challenge 4: Default-Action Complexity

Problem: The default action for many signals (kill, stop, ignore, or no-op) has grown into two nearly identical switch statements in `do_signal()`, once for the `SIG_DFL` path and again for handler-failure fallback. **Solution:** We consolidated the default-case logic into a single inline switch block that is invoked both when `sa_sigaction == SIG_DFL` and when our inline setup fails, reducing duplication and ensuring consistent behavior for each signal number.

4.5 Challenge 5: SIGALARM Implementation Complexity

Problem: Many files needed to be modified to support the optional function of an alarm, which added layers of difficulty to a project where most of the code is modified within a single file. **Solution:** Understanding the codebase using LLMs sped up the workflow and understanding of minor files such as `usys.pl` and `proc.h`.

5 Testing and Validation

We used the provided `basic.c` test suite—including `basic1--8`, `basic10--11`, and `basic20`—along with our own test cases: `siginfo`, `sigchld`, `alarm1`, and `alarm2`, to validate our implementation from Checkpoint 2 onward.

All tests passed, confirming correct behavior for signal delivery, handler installation, context switching, masking, timer signals, and metadata propagation.

5.1 Test Results

```
sh >> signal
=== TESTSUITE ===
- Project: signal test suite
Usage: ./signal [testname]
- [testname] can be one of the following:
  basic1
  basic2
  basic3
  basic4
  basic5
  basic6
  basic7
  basic10
  basic11
  basic20
  siginfo
  sigchld
  alarm1
  alarm2
  alarm3

Running all tests
signaltests starting
test basic1: OK
test basic2: OK
test basic3: OK
handler4 triggered
test basic4: OK
handler5 triggered
handler5 triggered
handler5 triggered
handler5 triggered
handler5 triggered
handler5 triggered
test basic5: OK
handler6 triggered due to 1
handler6_2 triggered due to 2
test basic6: OK
```

Figure 1: Final Test Results: Part1

```
handler7 triggered due to 1
handler7_2 triggered due to 2
test basic7: OK
test basic10: OK
test basic11: OK
test basic20: OK
test siginfo: OK
test sigchld: OK
alarm(2) triggered, 2 second timer starts
Alarm triggered!
test alarm1: OK
Starting test_alarm_remaining
alarm(5) returns remaining = 0
Waiting 2 seconds (busy-wait)...
Finished waiting 2 seconds.
alarm(3) returns remaining = 3
Waiting for alarm to trigger...
Alarm triggered!
alarm(5) returns remaining = 0
Waiting 2 seconds (busy-wait)...
Finished waiting 2 seconds.
alarm(0) returns remaining = 3
Waiting 5 seconds to confirm no alarm...
test_alarm_remaining completed
test alarm2: OK
alarm(10) triggered, 10 second timer starts
Waiting 2 seconds (busy-wait)...
Finished waiting 2 seconds.
Alarm cancelled early, remaining time: 8
test alarm3: OK
sh > child 4 exit with code 0
```

Figure 2: Final Test Results: Part2

Our implementation successfully passes all base checkpoint tests, demonstrating:

- Correct signal handler invocation
- Proper signal masking
- Signal delivery across process boundaries
- Context restoration functionality

6 Conclusion

We successfully completed Checkpoints 2 and beyond of the Signal project, building on the foundation from Checkpoint 1. Our implementation includes:

- **Checkpoint 2 (SIGKILL):** Added support for an uncatchable, unignorable SIGKILL that forcefully terminates a process.
- **Checkpoint 3 (fork/exec):** Implemented correct signal behavior across `fork` and `exec`.

- **Optional - SIGALRM:** Designed a user-mode timer syscall to raise **SIGALRM** after a specified duration.
- **Optional - SIGINFO:** Populated `siginfo_t` with detailed metadata passed to user handlers.
- **Optional - SIGCHLD:** Enabled parent processes to receive **SIGCHLD** with exit information upon child termination.

All functionalities were verified with custom test cases, ensuring correct behavior and kernel integration.